

The cdb-reading library interface

Overview

You can read records in a [constant database](#) from file descriptor fd as follows:

1. Use `cdb_init` to place information about fd into a `struct cdb` variable c .
2. Carry out any number of searches, as described below.
3. Use `cdb_free` to remove any memory map that might have been reserved by `cdb_init`.

Each search works as follows:

1. Use `cdb_find` to search for a record under key k . If `cdb_find` returns 0, the database does not contain that key; stop. If `cdb_find` returns -1, there was a read error; abort.
2. Use `cdb_datalen` to find the number of bytes of data in this record. Allocate a pointer d to a region of memory large enough to hold the data. If not enough memory is available, abort.
3. Use `cdb_read` with `cdb_datapos` to read the data. If `cdb_read` returns -1, there was a read error; abort.
4. Do something with the data, and then free the allocated region of memory.

There may be several records under a single key. You can use `cdb_findnext` to find the next record under this key.

Details

```
#include <cdb.h>

cdb_init(&c, fd);
cdb_free(&c);
result = cdb_read(&c, d, dlen, dpos);

cdb_findstart(&c);
result = cdb_findnext(&c, k, klen);
result = cdb_find(&c, k, klen);

dpos = cdb_datapos(&c);
dlen = cdb_datalen(&c);

static struct cdb c;
int fd;

char *d;
unsigned int dlen;
uint32 dpos;

char *k;
unsigned int klen;
int result;
```

A `struct cdb` variable such as c is either unallocated or allocated. If it is allocated, it holds information about a constant database:

- a file descriptor fd reading from the database;

- if convenient, a shared memory map reading from the database; and
- information about a search in progress.

`c` must be initialized to zero, meaning unallocated.

`cdb_free` unallocates `c` if `c` is allocated. Otherwise it leaves `c` alone. `cdb_free` does not close `fd`.

`cdb_init` allocates `c` to hold information about a constant database read by descriptor `fd`. You may call `cdb_init` repeatedly; if `c` is already allocated, `cdb_init` unallocates it first.

`cdb_read` reads `dlen` bytes into `d` from byte position `dpos` in the database. You must allocate `c` before calling `cdb_read`. Normally `cdb_read` returns 0. If the database file is shorter than `dpos+dlen` bytes, or if there is a disk read error, `cdb_read` returns -1, setting `errno` appropriately.

`cdb_findstart` prepares `c` to search for the first record under a new key. You must allocate `c` before calling `cdb_findstart`, and you must call `cdb_findstart` before calling `cdb_findnext`.

`cdb_findnext` looks for the n th record under key k in the database, where n is the number of calls to `cdb_findnext` after the most recent call to `cdb_findstart`. If it finds the record, `cdb_findnext` returns 1; if there are exactly $n-1$ such records, `cdb_findnext` returns 0; if there are fewer than $n-1$ such records, the behavior of `cdb_findnext` is undefined; if there is a database format error or disk error, `cdb_findnext` returns -1, setting `errno` appropriately. Each call to `cdb_findnext` (before another call to `cdb_findstart`) must use the same k and $klen$.

If `cdb_findnext` returns 1, it arranges for `cdb_datapos` to return the starting byte position of the data in the record, and for `cdb_datalen` to return the number of bytes of data in the record. Otherwise the results of `cdb_datapos` and `cdb_datalen` are undefined.

`cdb_find` is the same as `cdb_findstart` followed by `cdb_findnext`: it finds the first record under key k .

Beware that these functions may rely on non-atomic operations on the `fd` ofile, such as seeking to a particular position and then reading. Do not attempt two simultaneous database reads using a single ofile.