

# TinyCDB - a Constant DataBase

by Michael Tokarev, [mjt+cdb {at} tls {dot} msk {dot} ru](mailto:mjt+cdb@tls.msk.ru).

## Quick links

- [Introduction](#)
  - [Programming interface](#)
    - [Creating CDB file](#)
    - [Querying CDB file](#)
  - [Format of CDB file](#)
  - [Terms of usage](#)
  - [Download](#)
- 

## Introduction

TinyCDB is a very fast and simple package for creating and reading constant data bases, a data structure introduced by [Dan J. Bernstein](#) in his [cdb](#) package. It may be used to speed up searches in a sequence of (key,value) pairs with very big number of records. Example usage is indexing a big list of users - where a search will require linear reading of a large /etc/passwd file, and for many other tasks. It's usage/API is similar to ones found in [BerkeleyDB](#), [gdbm](#) and traditional \*nix dbm/ndbm libraries, and is compatible in great extent to `cdb-0.75` package by Dan Bernstein.

CDB is a *constant* database, that is, it cannot be updated at a runtime, only rebuilt. Rebuilding is atomic operation and is very fast - much faster than of many other similar packages. Once created, CDB may be queried, and a query takes very little time to complete.

## Programming interface

There are two interfaces provided by a library, `-lcdb`, -- create interface which is used to create CDB file, and two variants of query interface. A program using any routines should `#include <cdb.h>` header file which holds all required definitions of a library. More information together with detailed description of every routine is available in manual page inside TinyCDB package.

TinyCDB is different from Dan's `cdb-0.75` in the following ways:

- Only pure object-oriented approach is implemented, i.e. there is no documented way to directly write to or read from a CDB file. I found this approach more clean. With this, a record should fit in memory, but I don't know an application which may have a problem with that. All other hashing packages works the same way.
- Two query interfaces are provided - one compatible with Dan's `cdb-0.75`, and another which is compatible with older releases of his `cdb` library (which in fact allows direct reading of CDB file).
- At a time of CDB creation (`cdb_make`), it is possible to determine whenever a given key already exists (but this slows operation down, sometimes significantly), and to replace a key with a new one.
- TinyCDB usually requires much less memory to create CDB file, and never more than `cdb-0.75`.
- There is only one library, `-lcdb`, which required for linking and incorporates all the functionality (no additional libraries are required), and only one header file, `<cdb.h>`, which includes definitions for both `cdb` and `cdb_make` (it may be linked with `cdb_make.h` for programs that expects this header file).
- There is only one utility, `cdb`, which may be used for all the tasks including create, query, dump and so on. It understands Dan `cdbmake`'s input format, and traditional format understood by e.g. `makedb` utility found in BerkeleyDB package.

## Create interface

Create interface is built around `struct cdb_make` structure which is opaque type. The following is a sequence of action which should be performed in order to create CDB file (error handling is omitted):

```
struct cdb_make cdbm;
int fd;
char *key, *val;
unsigned klen, vlen;

fd = open(tmpfile, O_RDWR|O_CREAT); /* open temporary file */
cdb_make_start(&cdbm, fd); /* initialize structure */

cdb_make_add(&cdbm, key, klen, val, vlen)
/* add as many records as needed */

cdb_make_put(&cdbm, key, klen, val, vlen, flag);
/* alternative interface. flags is one of:
   CDB_PUT_ADD adds new record unconditionally like cdb_make_add()
   CDB_PUT_REPLACE if a key is already exists, replace the record
   CDB_PUT_INSERT add a record only if the key isn't already exists
*/

cdb_make_exists(&cdbm, key, klen);
/* a routine to test whenever a given key is already exists */

cdb_make_finish(&cdbm);
/* final stage - write indexes to CDB file */

rename(tmpfile, cdbfile);
/* atomically replace CDB file with newly built one */
```

## Query interface

There are two variants of query interface, one as found in `cdb-0.75`, and another as found in earlier versions of `cdb` (`cdb-0.6x`).

### Query interface 1

This interface is built around `struct cdb` structure which is opaque to the application. This interface designed to be efficient for many queries, for a single query second variant may be more efficient. The following is a sequence of calls needed to perform a query of a value in a CDB file:

```
int fd;
struct cdb cdb;
char *key, *val;
unsigned klen, vlen, vpos;

fd = open(cdbfile, O_RDONLY);
cdb_init(&cdb, fd); /* initialize internal structure */
if (cdb_find(&cdb, key, klen) > 0) { /* if search successeful */
    vpos = cdb_datapos(&cdb); /* position of data in a file */
    vlen = cdb_datalen(&cdb); /* length of data */
    val = malloc(vlen); /* allocate memory */
    cdb_read(&cdb, val, vlen, vpos); /* read the value into buffer */
    ... /* handle the value */
}
```

and here is what is needed to enumerate all values assotiated with a given key:

```
struct cdb_find cdbf; /* structure to hold current find position */
cdb_findinit(&cdbf, &cdb, key, klen); /* initialize search of key */
while(cdb_findnext(&cdbf) > 0) {
```

```

    vpos = cdb_datapos(&cdb);
    vlen = cdb_datalen(&cdb);
    val = malloc(vlen);
    cdb_read(&cdb, val, vlen, vpos);
    /* handle the value */
    free(val);
}

```

## Query interface 2

Another, simpler query interface exists which is suitable for a single query. Two routines provided works with a single filedescriptor opened for reading:

```

int fd;
char *key, *val;
unsigned klen
cdbi_t vlen;

fd = open(cdbfile, O_RDONLY); /* open a CDB file */
if (cdb_seek(fd, key, klen, &vlen) > 0) {
    /* if key was found, file will be positioned to the
     * start of data value and it's length will be placed to vlen */
    val = malloc(vlen);
    cdb_bread(fd, val, len); /* read the value;
                               * plain read() will do as well. */
    /* handle the value */
}

```

## Format of CDB file

To be written. Meanwhile, consult Dan Bernstein's [cdb manual](#).

## Terms of usage

The code is in public domain, that is, you may do anything you want with it.

## Download

Latest version is 0.78, released 11 May 2012, and can be found [here](#). It can be built on systems using RedHat Package Manager (rpm) with -tb option to create installable .rpm package. On a Debian GNU/Linux system, the preferred way to install it is to use standard apt repository. For other versions of the package and pre-built rpms look [here](#).

---

Enjoy. Michael Tokarev, mjt+cdb {at} tls {dot} msk {dot} ru.